

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
In [2]: df = pd.read_csv('cars24-car-price-clean.csv')
```

```
In [ ]: df.shape
```

```
In [ ]: df.head()
```

```
In [ ]: # Uni-variate
X = df['max_power'].values

# target
Y = df['selling_price'].values
```

```
In [ ]: X.shape
```

```
In [ ]: X.ndim
```

```
In [ ]: X = X.reshape(-1,1)
```

```
In [ ]: X.shape
```

```
In [ ]: X.ndim
```

```
In [ ]: Y.ndim
```

```
In [ ]: from sklearn.linear_model import LinearRegression
```

```
In [ ]: model = LinearRegression()
```

```
In [ ]: # this is where the entire training happens
model.fit(X, Y)
```

```
In [ ]: model.coef_
```

```
In [ ]: model.intercept_
```

```
In [ ]: type(model)
```

```
In [ ]: x_query = np.array([1.5])
x_query
```

```
In [ ]: model.predict(x_query.reshape(-1, 1))
```

```
In [ ]: y_pred = model.predict(X)
y_pred[:5]
```

```
In [ ]: Y[:5]
```

```
In [ ]: model.score(X, Y)
```

## Multiple Linear Regression

```
In [3]: df = pd.read_csv('cars24-car-price-clean.csv')
```

```
In [4]: df.head()
```

```
Out[4]:
```

	selling_price	year	km_driven	mileage	engine	max_power	age	make	model	Individual	Trustmark Dealer	Diesel	Electric	LPG
0	-1.111046	-0.801317	1.195828	0.045745	-1.310754	-1.157780	0.801317	-0.433854	-1.125683	1.248892	-0.098382	-0.985275	-0.020095	-0.056917
1	-0.223944	0.450030	-0.737872	-0.140402	-0.537456	-0.360203	-0.450030	-0.327501	-0.333227	1.248892	-0.098382	-0.985275	-0.020095	-0.056917
2	-0.915058	-1.426990	0.035608	-0.582501	-0.537456	-0.404885	1.426990	-0.327501	-0.789807	1.248892	-0.098382	-0.985275	-0.020095	-0.056917
3	-0.892365	-0.801317	-0.409143	0.329620	-0.921213	-0.693085	0.801317	-0.433854	-0.905265	1.248892	-0.098382	-0.985275	-0.020095	-0.056917
4	-0.182683	0.137194	-0.544502	0.760085	0.042999	0.010435	-0.137194	-0.246579	-0.013096	-0.800710	-0.098382	1.014945	-0.020095	-0.056917

```
In [5]: X = df.drop('selling_price', axis=1).values
```

```
In [6]: X.shape
```

```
Out[6]: (19820, 17)
```

```
In [7]: ones = np.ones((19820, 1))
X_new = np.hstack((ones,X))
X_new.shape
```

```
Out[7]: (19820, 18)
```

```
In [8]: X_new[:2]
```

```
Out[8]: array([[ 1.          , -0.80131654,  1.19582817,  0.04574517, -1.31075443,
                -1.15777962,  0.80131654, -0.43385435, -1.12568266,  1.24889206,
                -0.09838223, -0.9852749 , -0.02009467, -0.0569168 ,  1.0246219 ,
                0.4958182 ,  0.44450319, -0.42472845],
               [ 1.          ,  0.45003028, -0.73787208, -0.14040198, -0.53745638,
                -0.36020313, -0.45003028, -0.32750073, -0.3332271 ,  1.24889206,
                -0.09838223, -0.9852749 , -0.02009467, -0.0569168 ,  1.0246219 ,
                0.4958182 ,  0.44450319, -0.42472845]])
```

```
In [9]: Y = df['selling_price'].values
Y = Y.reshape(-1, 1)
Y.shape
```

```
Out[9]: (19820, 1)
```

```
In [ ]:
```

```
In [10]: def hypothesis(X, weights):
'''
    X : (n, d+1)
    weights : (d+1, 1)
'''
    return np.dot(X, weights)
```

```
In [11]: def error(X, Y, weights):
'''
    X : (n, d+1)
    Y : (n, 1)
    weights : (d+1, 1)
'''
    Y_hat = hypothesis(X, weights)
    err = np.mean((Y - Y_hat)**2)
    return err
```

```
In [12]: def gradients(X, Y, weights):
    Y_hat = hypothesis(X, weights)
    grads = np.dot( X.T , (Y_hat - Y ) )
    return 2*grads/len(Y)
```

```
In [16]: def gradient_descent(X, Y, max_itr = 200, learning_rate = 0.01):
```

```
    # step 1 : init() randomly
    weights = np.zeros((X.shape[1], 1))
    error_list = []

    # step 2 repeate until convergence
    for i in range(max_itr):

        e = error(X, Y, weights)
        error_list.append(e)

        grads = gradients(X, Y, weights)


        weights = weights - learning_rate*grads

    return weights, error_list
```

```
In [17]: opt_weights, error_list = gradient_descent(X_new, Y)
```

```
In [18]: plt.plot(error_list)
```

```
Out[18]: [<matplotlib.lines.Line2D at 0x7fd4b0d1ff70>]
```



```
In [39]: opt_weights.round(2)
```

```
Out[39]: array([[ 0.  ],
                [ 0.11],
                [-0.03],
                [-0.05],
                [ 0.07],
                [ 0.13],
                [-0.11],
                [ 0.18],
                [ 0.47],
                [-0.02],
                [-0.01],
                [ 0.04],
                [ 0.02],
                [ 0.01],
                [-0.03],
                [-0.07],
                [-0.  ],
                [ 0.  ]])
```

```
In [42]: # feature importances.  
np.abs(opt_weights.round(2))
```

```
Out[42]: array([[0.  ],  
               [0.11],  
               [0.03],  
               [0.05],  
               [0.07],  
               [0.13],  
               [0.11],  
               [0.18],  
               [0.47],  
               [0.02],  
               [0.01],  
               [0.04],  
               [0.02],  
               [0.01],  
               [0.03],  
               [0.07],  
               [0.  ],  
               [0.  ]])
```

```
In [22]: y_pred = hypothesis(X_new, opt_weights)
```

```
In [25]: y_pred[:5]
```

```
Out[25]: array([[ -1.20007006],  
               [-0.3020007  ],  
               [-0.95360925],  
               [-0.98039799],  
               [ 0.01098332]])
```

```
In [26]: Y[:5]
```

```
Out[26]: array([[ -1.11104589],  
               [-0.22394353],  
               [-0.91505816],  
               [-0.89236484],  
               [-0.18268296]])
```

```
In [29]: def r2_score(Y, Y_hat):  
          num = np.sum((Y - Y_hat)**2)  
          denom = np.sum((Y - Y.mean() )**2)  
          return (1 - num/denom)
```

```
In [30]: r2_score(Y, y_pred)
```

```
Out[30]: 0.9240707427695862
```

```
In [ ]:
```

```
In [37]: n = X.shape[0]  
          d = X.shape[1]  
  
          adj_r2 = 1 - ((1 - r2_score(Y, y_pred))*(n-1) / (n-d-1))  
          print(adj_r2)  
  
0.9240055575674391
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

In [ ]: